

kaggle房价预测笔记

STEP1:检视数据

```
import pandas as pd
import numpy as np

train_df = pd.read_csv("input/train.csv", index_col=0)
test_df = pd.read_csv("input/test.csv", index_col=0)

# 默认看前五行的
print(train_df.head())
```

STEP2:数据预处理

列表中很多数据都是英文单词还有些是无意义的数字，我们需要进行处理：

合并数据


我们要对所有数据进行处理，所以训练集和测试集都需要处理。

测试集中没有“SalePrice”列，所以我们需要对训练集里面的“SalePrice”进行处理工作。

```
# 观察SalePrice数据
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

train_df = pd.read_csv("input/train.csv", index_col=0)
test_df = pd.read_csv("input/test.csv", index_col=0)

prices = pd.DataFrame({"price": train_df["SalePrice"], "log(price+1)":
    np.log1p(train_df["SalePrice"])}))
prices.hist()
plt.show()
```



数据拼接

```
y_train = np.log1p(train_df.pop("SalePrice"))  
all_df = pd.concat((train_df, test_df), axis=0)
```

类特征工程处理

观察到“MSSubClass”类虽然都是数字表示，但是数字只是一个代号，并没有大小关系：

```

Id, MSSubClass, MSZoning, LotFrontage, LotArea, Street, Alley, LotShape, Lar
1, 60, RL, 65, 8450, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, CollgCr, Norm, Norm,
2, 20, RL, 80, 9600, Pave, NA, Reg, Lvl, AllPub, FR2, Gtl, Veenker, Feedr, Norm, 1F
3, 60, RL, 68, 11250, Pave, NA, IR1, Lvl, AllPub, Inside, Gtl, CollgCr, Norm, Norm
4, 70, RL, 60, 9550, Pave, NA, IR1, Lvl, AllPub, Corner, Gtl, Crawfor, Norm, Norm,
5, 60, RL, 84, 14260, Pave, NA, IR1, Lvl, AllPub, FR2, Gtl, NoRidge, Norm, Norm, 1F
6, 50, RL, 85, 14115, Pave, NA, IR1, Lvl, AllPub, Inside, Gtl, Mitchel, Norm, Norm
7, 20, RL, 75, 10084, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, Somerst, Norm, Norm
8, 60, RL, NA, 10382, Pave, NA, IR1, Lvl, AllPub, Corner, Gtl, NWAmes, PosN, Norm,
9, 50, RM, 51, 6120, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, OldTown, Artery, Nor
10, 190, RL, 50, 7420, Pave, NA, Reg, Lvl, AllPub, Corner, Gtl, BrkSide, Artery, A
11, 20, RL, 70, 11200, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, Sawyer, Norm, Norm
12, 60, RL, 85, 11924, Pave, NA, IR1, Lvl, AllPub, Inside, Gtl, NridgHt, Norm, Nor
13, 20, RL, NA, 12968, Pave, NA, IR2, Lvl, AllPub, Inside, Gtl, Sawyer, Norm, Norm
14, 20, RL, 91, 10652, Pave, NA, IR1, Lvl, AllPub, Inside, Gtl, CollgCr, Norm, Nor
15, 20, RL, NA, 10920, Pave, NA, IR1, Lvl, AllPub, Corner, Gtl, NAmes, Norm, Norm,
16, 45, RM, 51, 6120, Pave, NA, Reg, Lvl, AllPub, Corner, Gtl, BrkSide, Norm, Norm
17, 20, RL, NA, 11241, Pave, NA, IR1, Lvl, AllPub, CulDSac, Gtl, NAmes, Norm, Norm
18, 90, RL, 72, 10791, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, Sawyer, Norm, Norm
19, 20, RL, 66, 13695, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, SawyerW, RRAe, Nor
20, 20, RL, 70, 7560, Pave, NA, Reg, Lvl, AllPub, Inside, Gtl, NAmes, Norm, Norm, 1

```

对“MSSubClass”类的解释如下;

MSSubClass: Identifies the type of dwelling involved in the sale.

```

20  1-STORY 1946 & NEWER ALL STYLES
30  1-STORY 1945 & OLDER
40  1-STORY W/FINISHED ATTIC ALL AGES
45  1-1/2 STORY - UNFINISHED ALL AGES
50  1-1/2 STORY FINISHED ALL AGES
60  2-STORY 1946 & NEWER
70  2-STORY 1945 & OLDER
75  2-1/2 STORY ALL AGES
80  SPLIT OR MULTI-LEVEL
85  SPLIT FOYER
90  DUPLEX - ALL STYLES AND AGES
120 1-STORY PUD (Planned Unit Development) - 1946 & NEWER
150 1-1/2 STORY PUD - ALL AGES
160 2-STORY PUD - 1946 & NEWER
180 PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
190 2 FAMILY CONVERSION - ALL STYLES AND AGES

```

把category的变量转变成numerical表达形式

当我们用numerical来表达categorical的时候，要注意，数字本身有大小的含义，所以乱用数字会给之后的模型学习带来麻烦。于是我们可以用One-Hot的方法来表

pandas自带的get_dummies方法，可以帮你一键做到One-Hot。

在数据分析和机器学习领域中，"category"（或 "categorical"）变量是指具有有限个离散取值的变量，它描述了数据的类别或标签。

Category 变量通常用于表示具有固定类别的特征。这些类别可以是预定义的，例如性别（男、女）、颜色（红、绿、蓝）等，也可以是基于观察数据中的不同值进行自动识别的，例如用户的职业、产品的类型等。

与连续变量不同，Category 变量没有顺序或大小的概念，它们只是表示不同的类别或标签。在统计分析和机器学习中，处理 Category 变量通常需要进行特殊的编码或转换，以便将其表示为适合模型处理的数值形式。

常见的 Category 变量编码方法包括：

1. 无序编码（Nominal Encoding）：将每个类别分配一个唯一的整数或字符串标识符，例如使用独热编码（One-Hot Encoding）或标签编码（Label Encoding）。
2. 有序编码（Ordinal Encoding）：对类别进行排序，并为每个类别分配一个整数或字符串标识符，以表示它们之间的顺序关系。

Category 变量在数据分析和机器学习中扮演重要角色，因为它们可以提供有关数据的分类信息，帮助描述和解释数据，并在许多任务中用作特征输入。

独热编码

```
pd.get_dummies(all_df["MSSubClass"], prefix='MSSubClass')
```

以此类推，将所有的Category 变量使用One-Hot Encoding表示：

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

train_df = pd.read_csv("input/train.csv", index_col=0)
test_df = pd.read_csv("input/test.csv", index_col=0)

y_train = np.log1p(train_df.pop("SalePrice"))
all_df = pd.concat((train_df, test_df), axis=0)
all_dummy_df = pd.get_dummies(all_df["MSSubClass"], prefix='MSSubClass')
all_df = pd.concat([all_df, all_dummy_df], axis=1)
all_dummy_df = pd.get_dummies(all_df)
all_dummy_df.pop("MSSubClass")
print(all_dummy_df.head())
```

处理缺失数据

最常见的为直接删除，取平均值，取众数：

```
# 取平均值

mean_cols = all_dummy_df.mean()
all_dummy_df = all_dummy_df.fillna(mean_cols)
print(all_dummy_df.isnull().sum().sum())
```

标准化numerical数据

一般来说回归的分类器对源数据要求比较高，需要数据在标准分布里面，不要让数据间相差太大。

注意：OneHot数据不需要标准化

寻找需要标准化的数据：

```
# 所有数据类的的数据都需要标准化，
all_df = pd.concat((train_df, test_df), axis=0)
numeric_cols = all_df.columns[all_df.dtypes != 'object']
numeric_cols = numeric_cols.drop("MSSubClass")
```

标准化数据：

```
numeric_cols_means = all_dummy_df.loc[:, numeric_cols].mean()
numeric_cols_std = all_dummy_df.loc[:, numeric_cols].std()
all_dummy_df.loc[:, numeric_cols] = (all_dummy_df.loc[:, numeric_cols] -
    numeric_cols_means)/numeric_cols_std
```

重新获得训练集和测试集：

```
dummy_train_df = all_dummy_df.loc[train_df.index]
dummy_test_df = all_dummy_df.loc[test_df.index]
```

Step3 建立模型

代入模型寻找最优值：

```

alphas = np.linspace(0, 20, 20)
test_scores = []
for alpha in alphas:
    clf = Ridge(alpha)
    test_score = np.sqrt(-cross_val_score(clf, x_train, y_train, cv=10,
scoring='neg_mean_squared_error'))
    test_scores.append(np.mean(test_score))

```

```

# 使用 alpha=2.5 训练最终的 Ridge 模型
final_alpha = 2.5
final_clf = Ridge(final_alpha)
final_clf.fit(x_train, y_train)

# 预测测试集的 SalePrice
predicted_prices = np.expm1(final_clf.predict(x_test))

```


使用随机森林方法再做一次:

```


max_features = [.1, .3, .5, .7, .9, .99]
test_scores = []
for max_feat in max_features:
    clf = RandomForestRegressor(n_estimators=200, max_features=max_feat)
    test_score = np.sqrt(-cross_val_score(clf, x_train, y_train, cv=5,
scoring="neg_mean_squared_error"))
    test_scores.append(np.mean(test_score))

plt.plot(max_features, test_scores)
plt.title("max_features vs CV error")
plt.show()

```



进一步细分：



Ensemble

这里我们选择最简单的，取平均法：

```
# 使用 alpha=2.5 训练最终的 Ridge 模型
final_alpha = 2.5
final_clf = Ridge(final_alpha)
final_clf.fit(x_train, y_train)

# 预测测试集的 SalePrice
predicted_prices1 = np.expm1(final_clf.predict(x_test))

# 使用最佳参数模型
final_feature = 0.37
final_clf = RandomForestRegressor(n_estimators=200, max_features=final_feature)
final_clf.fit(x_train, y_train)

# 预测测试集的 SalePrice
predicted_prices2 = np.expm1(final_clf.predict(x_test))


# 将预测结果与对应的 Id 组合成 DataFrame
predictions_df = pd.DataFrame({"Id": test_df.index, "SalePrice":
(predicted_prices1+predicted_prices2)/2})

# 将 DataFrame 写入到 CSV 文件
predictions_df.to_csv("output/predictions_aver.csv", index=False)
```

使用BaggingRegressor对Ridge改进：

```
# 使用 alpha=15 训练最终的 Ridge 模型
final_alpha = 15
ridge = Ridge(final_alpha)
params = [1, 10, 15, 20, 25, 30, 40]
test_scores = []
for param in params:
    clf = BaggingRegressor(n_estimators=param, estimator=ridge)
    test_score = np.sqrt(-cross_val_score(clf, x_train, y_train, cv=10,
scoring="neg_mean_squared_error"))
    test_scores.append(np.mean(test_score))

plt.plot(params, test_scores)
plt.title("param vs CV error")
plt.show()
```

最终代码:

```
# 使用 alpha=15 训练最终的 Ridge 模型
final_alpha = 15
ridge = Ridge(final_alpha)
clf = BaggingRegressor(n_estimators=30, estimator=ridge)
clf.fit(x_train, y_train)

# 预测测试集的 SalePrice
predicted_prices2 = np.exp(m1(clf.predict(x_test)))

# 将预测结果与对应的 Id 组合成 DataFrame
predictions_df = pd.DataFrame({"Id": test_df.index, "SalePrice":
predicted_prices2})

# 将 DataFrame 写入到 CSV 文件
predictions_df.to_csv("output/predictions_BaggingRegressor_Ridge.csv",
index=False)
```


AdaBoostRegressor

```

# 使用 alpha=15 训练最终的 Ridge 模型
final_alpha = 15
ridge = Ridge(final_alpha)
params = [1, 10, 15, 20, 25, 30, 40, 45, 50]
test_scores = []
for param in params:
    clf = AdaBoostRegressor(n_estimators=param, estimator=ridge)
    test_score = np.sqrt(-cross_val_score(clf, x_train, y_train, cv=10,
scoring="neg_mean_squared_error"))
    test_scores.append(np.mean(test_score))

plt.plot(params, test_scores)
plt.title("param vs CV error")
plt.show()

```



最强XGBoost

```

# 使用 alpha=15 训练最终的 Ridge 模型
final_alpha = 15
ridge = Ridge(final_alpha)
params = [10, 15, 20, 30, 35, 40]
test_scores = []
for param in params:
    clf = XGBRegressor(n_estimators=param)
    test_score = np.sqrt(-cross_val_score(clf, x_train, y_train, cv=10,
scoring="neg_mean_squared_error"))
    test_scores.append(np.mean(test_score))

plt.plot(params, test_scores)
plt.title("param vs CV error")
plt.show()

```

